

Jacob Dineen

AI, Reasoning & Cognition (ARC) Lab — ARIZONA STATE UNIVERSITY

jdineen@asu.edu | (480) 603-6994 | jacobdineen.com

RESEARCH INTERESTS

I work on reasoning and alignment in large language models, with related interests in multi-agent reinforcement learning, controllability, and explainable AI (XAI). I have also held data science, ML engineering, and applied research roles across financial services and AI startups.

EDUCATION

ARIZONA STATE UNIVERSITY — Ph.D. Artificial Intelligence GPA: 4.00/4.00 2022–2027

Advisor: Ben Zhou. Committee: Muhao Chen, Chitta Baral, Vivek Gupta.

UNIVERSITY OF VIRGINIA — M.Sc. Computer Science GPA: 3.96/4.00 2019–2021

Advisor: Madhav Marathe.

SYRACUSE UNIVERSITY — M.S. Data Science GPA: 4.00/4.00 2017–2018

GRAND CANYON UNIVERSITY — B.S. Finance & Economics GPA: 3.65/4.00 2012–2015

RESEARCH EXPERIENCE

Graduate Research Assistant, ARC Lab (ARIZONA STATE UNIVERSITY) 2024–Present

- Research on reasoning, alignment, controllability, and multi-agent RL in LLMs.

Graduate Research Assistant, SEFCOM (ARIZONA STATE UNIVERSITY) 2022–2024

- Research at the intersection of AI and cybersecurity.

Applied Research, CAPITAL ONE — Center for Machine Learning (C4ML) 2020–2021

- Built reinforcement learning and agent-based simulations for organizational dynamics.

Research Assistant, BIOCOMPLEXITY INSTITUTE (UNIVERSITY OF VIRGINIA) 2019–2020

- Worked on graph dynamical systems, cooperative game theory, and behavioral modeling.

PROFESSIONAL EXPERIENCE

AI/ML Engineer (Contractor), ALUMBRA AI 2026–Present

Research (Contractor), PARETO AI 2025–Present

Machine Learning Engineer, SPRING OAKS CAPITAL 2022–2025

Data Scientist, CAPITAL ONE 2021–2022

Ph.D. Internships (3×), CAPITAL ONE 2020–2021

Analyst & Business Intelligence, REAL WORLD MARKETING 2016–2019

Data Scientist, BUFFALO CHECK LLC

2015–2019

Optimization Analyst, VOLTARI

2012–2015

PUBLICATIONS

Google Scholar Profile

* Equal Contribution, + Corresponding Author / Mentor

Peer-Reviewed Conference Proceedings (C)

- C1. **Dineen, Jacob**⁺, RRV, Aswin, Liu, Qin, Xu, Zhikun, Ye, Xiao, Shen, Ming, Li, Zhaonan, Lu, Shijie, Baral, Chitta, Chen, Muhao, & Zhou, Ben (2025). *QA-LIGN: Aligning LLMs through Constitutionally Decomposed QA*. In *Findings of the Association for Computational Linguistics: EMNLP 2025*, pp. 20619–20642.
- C2. RRV, Aswin, **Dineen, Jacob**, Handa, Divij, Uddin, Md Nayem, Parmar, Mihir, Baral, Chitta, & Zhou, Ben (2025). *ThinkTuning: Instilling Cognitive Reflections without Distillation*. In *Proceedings of the 2025 Conference on Empirical Methods in Natural Language Processing*, pp. 31236–31250.
- C3. Xu, Zhikun, Shen, Ming, **Dineen, Jacob**, Li, Zhaonan, Ye, Xiao, Lu, Shijie, RRV, Aswin, Baral, Chitta, & Zhou, Ben (2025). *ToW: Thoughts of Words Improve Reasoning in Large Language Models*. In *Proceedings of the 2025 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies (Volume 1: Long Papers)*, pp. 3057–3075.
- C4. **Dineen, Jacob**⁺, Kridel, Donald J., Dolk, Daniel R., & Castillo, David G. (2023). *Unified Explanations in Machine Learning Models: A Perturbation Approach*. In *Proceedings of the 56th Hawaii International Conference on System Sciences (HICSS-56)*, pp. 795–804.
- C5. **Dineen, Jacob**⁺, Haque, A. S. M. Ahsan-Ul, & Bielskas, Matthew (2021a). *Formal Methods for an Iterated Volunteer's Dilemma*. In *Proceedings of the 14th International Conference on Social, Cultural, and Behavioral Modeling (SBP-BRiMS 2021)*, pp. 81–90.
- C6. **Dineen, Jacob**⁺, Haque, A. S. M. Ahsan-Ul, & Bielskas, Matthew (2021b). *Reinforcement Learning for Data Poisoning on Graph Neural Networks*. In *Proceedings of the 14th International Conference on Social, Cultural, and Behavioral Modeling (SBP-BRiMS 2021)*, pp. 141–150.
- C7. Dolk, Daniel R., Kridel, Donald J., **Dineen, Jacob**⁺, & Castillo, David G. (2020). *Model Interpretation and Explainability towards Creating Transparency in Prediction Models*. In *Proceedings of the 53rd Hawaii International Conference on System Sciences (HICSS-53)*.

Peer-Reviewed Workshop Papers (WS)

- WS1. Srinivasan, Adarsh, **Dineen, Jacob**⁺, Afzal, Muhammad Umar, Sarfraz, Muhammad Uzair, Riaz, Irbaz Bin, Zhou, Ben (2026). *RECAP: Transparent Inference-Time Emotion Alignment for Medical Dialogue Systems*. In *Proceedings of the 8th Clinical Natural Language Processing Workshop (Oral)*.
- WS2. Li, Zhaonan, Chickering, Kyle R., Li, Bangzheng, **Dineen, Jacob**, Ye, Xiao, Xu, Zhikun, Lu, Shijie, Huang, Yuxi, Shen, Ming, Nguyen, Bach, Pavuluri, Jaya Adithya, Nguyen, Mau Son, Chavan, Sanika, Le, Ngoc Minh Thu, Chen, Muhao, Zhou, Ben (2026). *VisAnalog: A Diagnostic Suite for Visual Concept Transfer on Natural Images*. CVPR Workshop on Visual Concepts (VisCon), 2026 (Oral).

Working Papers / Under Review (W)

- W1. Li, Zhaonan, Lu, Shijie, Wang, Fei, **Dineen, Jacob**, Ye, Xiao, Xu, Zhikun, Liu, Siyi, Cho, Young Min, Li, Bangzheng, Chang, Daniel, Nguyen, Kenny, Yang, Qizheng, Chen, Muhao, Zhou, Ben (2025). *Unbiased Visual Reasoning with Controlled Visual Inputs*. Under review at COLM 2026.
- W2. Liu, Qin, **Dineen, Jacob**, Huang, Yuxi, Zhang, Sheng, Poon, Hoifung, Zhou, Ben, Chen, Muhao (2025). *ArenaBench: Automatic Benchmark Evolution via Multi-Model Competitive Evaluation*. Under review at COLM 2026.
- W3. Ye, Xiao, Li, Zhaonan, **Dineen, Jacob**, Xu, Zhikun, Lu, Shijie, Shen, Ming, Shrivastava, Shaswat, Ahuja, Avneet, Zhou, Ben (2025). *CC-LEARN: Cohort-Based Consistency Learning*. Under review at COLM 2026.
- W4. Shen, Ming, Xu, Zhikun, **Dineen, Jacob**, Ye, Xiao, Zhou, Ben (2025). *Training Language Models with Context-Free Reasoning Bottlenecks*. Under review at COLM 2026.
- W5. RRV, Aswin, **Dineen, Jacob**, Handa, Divij, Parmar, Mihir, Zhou, Ben, Baral, Chitta, Mishra, Swaroop (2026). *Mid-training with Self-Generated Data Improves Reinforcement Learning in Language Models*. Under review at NeurIPS 2026.
- W6. **Dineen, Jacob**, RRV, Aswin, Xu, Zhikun, Zhou, Ben (2026). *Vocabulary Dropout for Curriculum Diversity in LLM Co-Evolution*. Under review at COLM 2026.
- W7. Xu, Zhikun, Feng, Yu, **Dineen, Jacob**, Zhou, Ben (2026). *Skill Reuse as Compression in Agentic RL*. Under review at EMNLP 2026.
- W8. Ye, Xiao, **Dineen, Jacob**, Li, Zhaonan, Xu, Zhikun, Chen, Weiyu, Lu, Shijie, Huang, Yuxi, Shen, Ming, Tran, Phu, Yum, Ji-Eun Irene, Khan, Muhammad Ali, Afzal, Muhammad Umar, Riaz, Irbaz Bin, Zhou, Ben (2025). *Evaluating Medical LLMs by Levels of Autonomy: A Survey*. In preparation.

SKILLS

OS: Linux (Ubuntu), macOS, Windows

Languages: Python, Rust, C, C++, Java, JavaScript, R, Bash, PRISM, x86-64

Databases: MySQL, PostgreSQL, MongoDB, Snowflake, Redis, SQLite, Redshift

Markup: $\LaTeX$, HTML

ML Libraries: PyTorch, TensorFlow, Keras, JAX, Numpy, Pandas, Polars, Dask, HuggingFace, TRL, vLLM, VeRL, PySpark, NetworkX, DGL, Torch Geometric, BoTorch, SnowparkML

Tools: Git, VSCode, Docker, Kubernetes, Helm, Airflow, AWS (ECR/S3/EKS/CodeBuild), Databricks, OR-Tools, Sigma, Streamlit

SERVICE

Conference Reviewer: HICSS, SBP-BRiMS, NAACL, EMNLP, EACL, COLM

Teaching / Tutoring:

- CGCC Calculus & Linear Algebra Tutor (2019)
- Teaching Assistant — ASU CSE 365 / pwn.college (Security)

MISCELLANEOUS

Expert AI Trainer, PARETO AI (2024–Present)

Independent Contractor, MERCOR INTELLIGENCE (2025–Present)

Cybersecurity: pwn.college green belt (binary exploitation), user: jdin

GRADUATE COURSEWORK

ARIZONA STATE UNIVERSITY — PhD Computer Science:

Knowledge Representation; Computer Systems Security; Software Security; Planning and Learning Methods in AI; Algorithms; Research/Dissertation Hours.

UNIVERSITY OF VIRGINIA — MSc Computer Science:

Algorithms; Machine Learning; Computer Vision; Formal Methods; Reinforcement Learning; Graph Mining; Learning Theory (Game Theory); Cloud Computing; Research Hours.

SYRACUSE UNIVERSITY — MS Data Science:

Data Analysis and Decision Making; Business Analytics; Financial Analytics; Marketing Analytics; Advanced Information Systems; Data Science; Data Warehousing; Text Mining; Scripting for Data Analysis; Information Policy.